MedResearchBench: A Multi-Domain Benchmark for Evaluating AI Research Agents on Clinical Medical Research

Zhanxiao Tian[1]

[1]Independent Researcher

Preprint | arXiv cs.AI

March 2026

Abstract

The rapid advancement of AI research automation systems—including AI Scientist, data-to-paper, and Agent Laboratory—has demonstrated the potential for autonomous scientific discovery. However, existing benchmarks for evaluating these systems focus predominantly on fundamental sciences (machine learning, physics, chemistry), overlooking the unique challenges of medical clinical research: complex survey designs, inferential statistics with confounding control, adherence to reporting standards (STROBE, CONSORT), and the requirement for clinically actionable interpretation. We present **MedResearchBench**, the first benchmark specifically designed to evaluate AI systems on medical clinical research tasks. MedResearchBench comprises 16 tasks spanning 8 clinical domains (cardiovascular, oncology, mental health, metabolic, respiratory, neurology, infectious disease), built on publicly available datasets (NHANES, SEER) with ground truth from 16 high-quality published papers (IF range: 2.3–51.0). Each task is evaluated along 6 medical-specific dimensions: statistical methodology, results accuracy, visualization quality, clinical interpretation, confounding sensitivity, and reporting compliance. We describe the benchmark design rationale, task construction methodology, paper selection criteria with anti-paper-mill filtering, and a detailed analysis of task characteristics including methodological diversity, evaluation dimension coverage, and difficulty stratification. To demonstrate benchmark executability, we evaluate an agentic data2paper pipeline on 3 pilot tasks spanning all three difficulty tiers, achieving scores of 72/100 (Tier 1, Cardio_000), 69/100 (Tier 2, Mental_000), and 75/100 (Tier 3, Metabolic_002), with a mean score of 72/100 (B-level). Survey-weighted methodology was correctly implemented across all tasks; primary limitations included covariate incompleteness and reference group misspecification. MedResearchBench addresses a critical gap in AI research evaluation and provides a standardized, community-extensible platform for assessing whether AI systems can conduct clinically sound, publication-quality medical research. All task materials are publicly available at https://github.com/TerryFYL/MedResearchBench.

Keywords: AI research automation, medical informatics, benchmark, clinical research, NHANES, SEER, research quality evaluation

---

Contents

1 Introduction **4**
1.1 The Rise of AI Research Automation . . . 4
1.2 The Evaluation Gap . . . 4
1.3 The NHANES Paper Mill Problem . . . 5
1.4 Contributions . . . 5

2 Related Work **6**
2.1 AI Research Automation Systems . . . 6
2.2 Research Quality Benchmarks . . . 6
2.3 Medical Research Quality Standards . . . 7

3 Benchmark Design **7**
3.1 Overview . . . 7
3.2 Domain Coverage . . . 7
3.3 Task Structure . . . 8
3.4 Ground Truth Paper Selection . . . 8
3.5 Evaluation Dimensions . . . 9
3.6 Scoring Methodology . . . 9
3.7 Study Design Coverage . . . 9

4 Task Analysis and Characteristics **9**
4.1 Methodological Diversity . . . 10
4.2 Evaluation Dimension Distribution . . . 10
4.3 Ground Truth Paper Quality Distribution . . . 11
4.4 Difficulty Stratification . . . 11
4.5 Dataset Utilization . . . 12
4.6 Comparison with ResearchClawBench Task Characteristics . . . 13
4.7 Initial Baseline Evaluation . . . 13

5 Discussion **14**
5.1 Why Medical Research Needs Its Own Benchmark . . . 14
5.2 Comparison with ResearchClawBench . . . 15
5.3 Responsible Use: Quality Gates Against Paper Mills . . . 15
5.4 Limitations . . . 15
5.5 Future Directions . . . 16

6 Conclusion **16**

A Complete Task List with Metadata **18**

B LLM Judge Prompt Template **19**

C Paper Selection Criteria **20**

1 Introduction

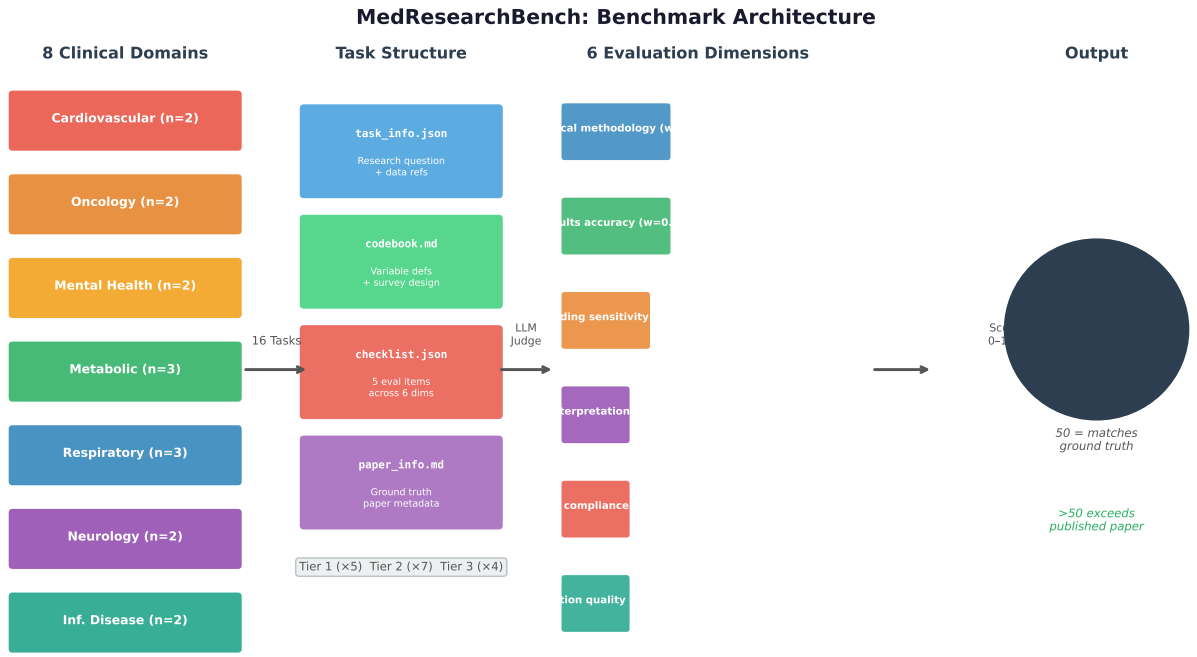

Figure 1: MedResearchBench architecture overview. Tasks are organized across 8 clinical domains and evaluated by an LLM Judge along 6 medical-specific dimensions. Scores range 0–100, where 50 indicates matching the published ground truth paper quality.

1.1 The Rise of AI Research Automation

Artificial intelligence systems capable of conducting scientific research autonomously have emerged as a transformative force in computational science. Lu et al. (2024) introduced *The AI Scientist*, a fully automated system capable of generating research ideas, conducting experiments, and writing complete papers, published in Nature. Independently, Ifargan et al. (2024) developed *data-to-paper*, a system that generates complete research papers from raw datasets, reported in NEJM AI. Schmidgall et al. (2025) presented *Agent Laboratory*, a multi-agent framework for autonomous research, accepted at EMNLP. These systems and others—including AI-Researcher (Huang et al., 2025, NeurIPS Spotlight), Virtual Lab (Swanson et al., 2024, Nature), and Robin (FutureHouse, 2025)—represent a rapidly growing ecosystem of automated research tools.

1.2 The Evaluation Gap

As these systems proliferate, the need for rigorous evaluation becomes paramount. ResearchClawBench (Schmidgall et al., 2025) represents the most comprehensive effort to date,

providing a benchmark of research tasks across 10 fundamental science domains with LLM-based evaluation. ScienceBench and related benchmarks focus on specific scientific reasoning capabilities.

However, **medical clinical research presents fundamentally different challenges** that existing benchmarks do not capture:

Table 1: Challenges distinguishing medical clinical research from fundamental science benchmarks.

Challenge	**Fundamental Science**	**Medical Clinical Research**
Data type	Experimental/simulation	Real-world population data (surveys, EHR, registries)
Statistical paradigm	Algorithm metrics (accuracy, F1)	Inferential statistics (OR, HR, regression, survival)
Confounding	Controlled experiments	Must be statistically controlled (DAG, propensity score)
Reporting standards	None formalized	STROBE, CONSORT, PRISMA (hard requirements)
Output requirement	"We found X"	"Clinicians should do Y because Z"
Sampling complexity	i.i.d. assumed	Complex survey design, stratification, weights
Ethics dimension	Minimal	Subgroup fairness, clinical harm potential

No existing benchmark tests whether an AI system can correctly go from patient-level data to a clinically sound, publication-ready manuscript.

1.3 The NHANES Paper Mill Problem

The urgency of quality evaluation in medical AI research is underscored by the NHANES paper mill scandal. In 2025, a PLOS Biology investigation identified hundreds of formulaic papers exploiting the publicly available National Health and Nutrition Examination Survey (NHANES) dataset, characterized by mechanical application of logistic regression without genuine clinical insight, inadequate confounding control, and failure to account for complex survey design. These papers, while technically "published," represent a degradation of the medical literature.

AI research automation systems risk amplifying this problem at unprecedented scale—or, if properly evaluated, providing the quality gates to prevent it. MedResearchBench is designed to distinguish between systems that produce rigorous, clinically meaningful research and those that merely generate publication-shaped text.

1.4 Contributions

We make the following contributions:

1. **MedResearchBench**: The first benchmark specifically designed for evaluating AI systems on medical clinical research tasks, comprising 16 tasks across 8 clinical domains.

2. **Medical-specific evaluation dimensions**: A 6-dimension scoring framework that captures the unique quality requirements of clinical research, including confounding sensitivity and reporting compliance.

3. **Anti-paper-mill design**: Explicit quality criteria that distinguish genuine research from formulaic output, directly addressing the NHANES paper mill problem.

4. **Public benchmark with community extensibility**: All tasks use publicly available datasets (NHANES, SEER), enabling reproducible evaluation and community contribution.

5. **Initial baseline results**: End-to-end evaluation of an agentic pipeline across 3 pilot tasks (Tier 1–3), yielding a mean score of 72/100 (B-level), establishing the first quantitative baseline for AI-driven medical research quality on this benchmark.

2 Related Work

2.1 AI Research Automation Systems

The landscape of AI research automation has evolved rapidly since 2024. We categorize existing systems by their scope and approach:

End-to-end systems attempt to automate the full research pipeline from idea generation to paper writing. The AI Scientist (Lu et al., 2024) pioneered this approach in machine learning, generating novel research ideas, implementing experiments, and producing complete LaTeX papers. data-to-paper (Ifargan et al., 2024) focuses on the data-to-manuscript pipeline, using LLM agents with rule-based guardrails to generate papers from raw datasets. Agent Laboratory (Schmidgall et al., 2025) employs a multi-agent architecture with specialized roles (literature review, experiment design, writing).

Domain-specific systems target particular research areas. Virtual Lab (Swanson et al., 2024) combines LLM agents with computational biology tools for nanobody design. Robin (FutureHouse, 2025) orchestrates web-scale literature retrieval for biomedical research. AI-Researcher (Huang et al., 2025) introduces autonomous research with open-ended exploration capabilities.

Medical research tools remain comparatively underdeveloped. While clinical decision support and literature review tools exist, no published system specifically targets the end-to-end automation of observational clinical research—the dominant paradigm in medical science.

2.2 Research Quality Benchmarks

ResearchClawBench (Schmidgall et al., 2025) is the most directly relevant benchmark, providing tasks across 10 scientific domains with a dual-mode LLM Judge evaluation. Tasks include data, related work references, and target study checklists. However, its domain coverage—machine learning, NLP, computer vision, reinforcement learning, social science, statistics, economics, environmental science, psychology, political science—excludes clinical medicine entirely.

ScienceBench and related efforts evaluate specific scientific reasoning capabilities (mathematical derivation, experimental design) but do not assess the integrated research pipeline from data to publication.

No existing benchmark evaluates AI systems on medical clinical research tasks.

2.3 Medical Research Quality Standards

Medical research quality is governed by well-established reporting guidelines:

- **STROBE** (Strengthening the Reporting of Observational Studies in Epidemiology): Checklist for cross-sectional, cohort, and case-control studies
- **CONSORT** (Consolidated Standards of Reporting Trials): Checklist for randomized controlled trials
- **PRISMA** (Preferred Reporting Items for Systematic Reviews and Meta-Analyses): Checklist for systematic reviews

These standards provide an objective framework for evaluating AI-generated medical research—a unique advantage over fundamental science domains where reporting standards are less formalized.

3 Benchmark Design

3.1 Overview

MedResearchBench comprises 16 tasks (Phase 1) spanning 8 clinical domains, designed to evaluate AI systems' ability to conduct complete medical clinical research from raw data to publication-quality output. Each task is anchored by a ground truth published paper that used publicly available data, enabling objective evaluation.

3.2 Domain Coverage

We selected 8 clinical domains based on three criteria: (1) public data availability, (2) clinical impact (global disease burden), and (3) methodological diversity.

Table 2: Domain coverage and primary methods in MedResearchBench.

Domain	**# Tasks**	**Primary Datasets**	**Key Methods**
Cardiovascular	2	NHANES	Survey-weighted regression, dose-response
Oncology	2	SEER, SEER-Medicare	Survival analysis, temporal trends
Mental Health	2	NHANES	Piecewise regression, trend analysis
Metabolic	3	NHANES, NHANES III	ROC analysis, Cox PH, RCS
Respiratory	3	NHANES	Biomarker-exposure regression, mortality linkage
Neurology	2	NHANES	Cognitive assessment, composite index
Infectious Disease	2	NHANES	Temporal trends, vaccine effectiveness

3.3 Task Structure

Each task follows a standardized directory structure compatible with ResearchClawBench:

```
tasks/{TaskID}/
|-- task_info.json       # Task description, data references,
   metadata
|-- data/
|   |-- codebook.md      # Variable definitions, survey design
   variables
|   `-- placeholder.txt  # Data source and download
   instructions
|-- related_work/
|   `-- references.md    # 3-4 related papers (title, DOI)
`-- target_study/
    |-- paper_info.md    # Ground truth paper metadata
    |-- checklist.json   # 5 evaluation items across 6
       dimensions
    `-- images/          # Figure description placeholders
```

The `task_info.json` file extends the ResearchClawBench format with medical-specific metadata. The `codebook.md` file is a medical-specific addition providing variable definitions essential for correct analysis of complex survey data, including survey design variables (weights, PSU, strata).

3.4 Ground Truth Paper Selection

For each task, we selected a published paper meeting the following criteria:

1. **Authenticity**: Paper existence verified via PubMed/Google Scholar (no fabricated references)
2. **Quality**: Published in a peer-reviewed journal with Impact Factor ≥ 2.0
3. **Relevance**: Uses the specified public dataset to address the specified research question
4. **Methodological rigor**: Employs appropriate statistical methods, adequate confounding control, and follows reporting standards
5. **Anti-paper-mill filtering**: Excludes formulaic single-factor studies without genuine analytical innovation

Ground truth papers span a wide quality range (IF 2.3–51.0), reflecting the realistic distribution of medical publications. This range is deliberate: AI systems should be evaluated against papers of varying quality tiers, not only elite publications. Publication years range from 2015 to 2026, with 13 papers (81%) published in 2020 or later, ensuring methodological currency.

Table 3: Ground truth papers for MedResearchBench Phase 1 tasks.

Task ID	Ground Truth Paper	Journal	IF
Cardio_000	Luo et al. (2026) Dietary sodium and hypertension	PLOS ONE	∼3.7
Cardio_003	Aggarwal et al. (2021) Racial disparities in HTN	Hypertension	∼8.3
Onco_000	Pankratz et al. (2022) CRC survival trends	Cancer Control	∼3.0
Onco_002	Liang et al. (2020) SES disparities in CRC	Clin Transl Gastroenterol	∼3.6
Mental_000	Dong et al. (2022) Sleep duration and depression	J Affect Disord	5.4
Mental_002	Yu et al. (2020) Depression prevalence trends	J Affect Disord	5.4
Metabolic_000	Guess et al. (2023) HOMA-IR threshold for MetS	Metab Syndr Relat Disord	∼2.3
Metabolic_002	Zhu et al. (2024) TSH and mortality	Cardiovasc Diabetol	∼8.5
Metabolic_003	Wang et al. (2024) NAFLD prevalence	Ann Hepatol	3.4
Neuro_002	Lu et al. (2023) Heavy metals and cognition	Br J Nutr	∼3.3
Neuro_003	Huang et al. (2025) Antioxidants and PD mortality	Front Aging Neurosci	∼4.1
Infection_000	Le et al. (2020) HBV seroprevalence trends	JAMA Network Open	∼13.0
Infection_001	Rosenblum et al. (2022) HPV vaccine effectiveness	Ann Intern Med	∼51.0
Respiratory_000	Mendy et al. (2022) VOC and lung function	Eur Respir J	∼16.0
Respiratory_001	Martinez et al. (2015) Undiagnosed COPD	Ann Am Thorac Soc	∼6.0
Respiratory_003	Yin et al. (2025) METS-IR and OSA risk	Sci Rep	∼3.8

3.5 Evaluation Dimensions

MedResearchBench evaluates AI output along 6 dimensions, each addressing a critical aspect of medical research quality:

Each task's `checklist.json` contains 5 evaluation items distributed across at least 4 of these 6 dimensions.

3.6 Scoring Methodology

We adopt a dual-mode LLM Judge evaluation approach compatible with ResearchClawBench:

Objective Mode (primary): Each checklist item is scored 0–100, where 50 represents matching the ground truth paper's quality. Scores above 50 indicate the AI output exceeded the published paper on that specific criterion. The weighted sum across all items produces the final task score.

Subjective Mode (supplementary): Open-ended assessment of overall research quality, novelty, and clinical utility.

Medical-specific scoring extensions include a *survey design penalty* (failure to account for complex survey design results in automatic score reduction for `statistical_methodology`), a *confounding cascade* (progressive model adjustment evaluation), and a *STROBE compliance checklist* mapping to specific items relevant to each study design.

3.7 Study Design Coverage

4 Task Analysis and Characteristics

To characterize the benchmark's coverage and difficulty landscape, we analyze the 16 Phase 1 tasks across multiple dimensions.

Table 4: Evaluation dimensions and their medical-specific rationale.

Dimension	Description	Why Medical-Specific
`statistical_methodology`	Appropriate method selection, model specification, assumption checking	Complex survey design requires specialized methods
`results_accuracy`	Correct numerical results, properly constructed tables/figures	Medical statistics have specific conventions (OR with 95% CI)
`visualization_quality`	Clear, informative, publication-ready visualizations	Domain-specific figure types (KM curves, forest plots)
`clinical_interpretation`	Clinically meaningful discussion, actionable conclusions	Medicine requires "clinicians should do Y"; not just "we found X"
`confounding_sensitivity`	Adequate identification and control of confounders	Defining challenge of observational research
`reporting_compliance`	Adherence to STROBE/CONSORT/PRISMA standards	Hard requirements in medical journals

Table 5: Study design distribution across MedResearchBench Phase 1 tasks.

Study Design	# Tasks	Examples
Cross-sectional	7	Cardio_000, Mental_000, Metabolic_000, Neuro_002
Serial/repeated cross-sectional (trend)	4	Infection_000, Infection_001, Mental_002, Metabolic_003
Prospective cohort (mortality linkage)	3	Metabolic_002, Neuro_003, Respiratory_001
Retrospective cohort (registry)	2	Onco_000, Onco_002

4.1 Methodological Diversity

Across 16 tasks, we identify 23 distinct statistical methods. Table 6 summarizes the most frequently required methods.

Survey-weighted regression is near-universal (14/16 tasks), creating a natural "baseline competency test": any AI system that cannot handle complex survey design will fail across the majority of tasks.

4.2 Evaluation Dimension Distribution

Four dimensions achieve universal coverage, reflecting their centrality to medical research quality. `visualization_quality` appears in only 2 tasks in Phase 1; Phase 2 will systematically increase visualization coverage.

Table 6: Statistical methods required across MedResearchBench tasks.

Method	# Tasks	Representative Tasks
Survey-weighted regression (logistic/linear)	14	All NHANES-based tasks
Subgroup/stratified analysis	12	Cardio_003, Neuro_002, Respiratory_003
Restricted cubic splines (RCS)	5	Cardio_000, Metabolic_002, Neuro_003
Cox proportional hazards	4	Metabolic_002, Neuro_003, Onco_002, Respiratory_001
Temporal trend analysis (P-for-trend)	4	Infection_000, Infection_001, Mental_002
Sensitivity analysis	4	Cardio_000, Respiratory_000, Metabolic_002, Neuro_003
Kaplan-Meier / survival estimation	2	Onco_000, Onco_002
ROC curve / diagnostic performance	1	Metabolic_000
Vaccine effectiveness estimation	1	Infection_001
Piecewise linear regression	1	Mental_000

Table 7: Evaluation dimension coverage across 16 tasks.

Dimension	# Tasks Covering	Coverage Rate
`statistical_methodology`	16/16	100%
`results_accuracy`	16/16	100%
`clinical_interpretation`	16/16	100%
`reporting_compliance`	16/16	100%
`confounding_sensitivity`	14/16	87.5%
`visualization_quality`	2/16	12.5%

4.3 Ground Truth Paper Quality Distribution

The 50:31:19 (C:B:A) ratio approximates the real-world quality pyramid of medical publications. AI systems should be evaluated not only against elite papers but also against the "typical" published research.

4.4 Difficulty Stratification

We propose a three-tier difficulty classification based on the number of distinct analytical components required:

Tier 1 — Foundational (5 tasks): Single study design, standard regression, straightforward outcome. Examples: Cardio_000 (cross-sectional logistic regression), Metabolic_003 (prevalence estimation).

Tier 2 — Intermediate (7 tasks): Multiple analytical stages, dose-response modeling, or temporal trend analysis requiring multi-cycle data harmonization. Examples: Mental_000 (U-shaped dose-response with threshold estimation), Infection_000 (9-cycle temporal trend + birth cohort analysis).

Tier 3 — Advanced (4 tasks): Hybrid study designs, mortality linkage, causal inference frameworks, or policy-period natural experiments. Examples: Metabolic_002 (prospective cohort with NDI-linked Cox PH + nonlinear RCS), Infection_001 (quasi-experimental vaccine

Table 8: Ground truth paper quality tier distribution.

Tier	IF Range	# Papers	Proportion	Journals
A (Elite)	≥10	3	19%	Ann Intern Med, JAMA Network Open, Eur Respir J
B (Strong)	5–10	5	31%	Hypertension, J Affect Disord, Cardiovasc Diabetol
C (Solid)	2–5	8	50%	PLOS ONE, Cancer Control, Br J Nutr, Sci Rep, ...

impact + herd immunity quantification).

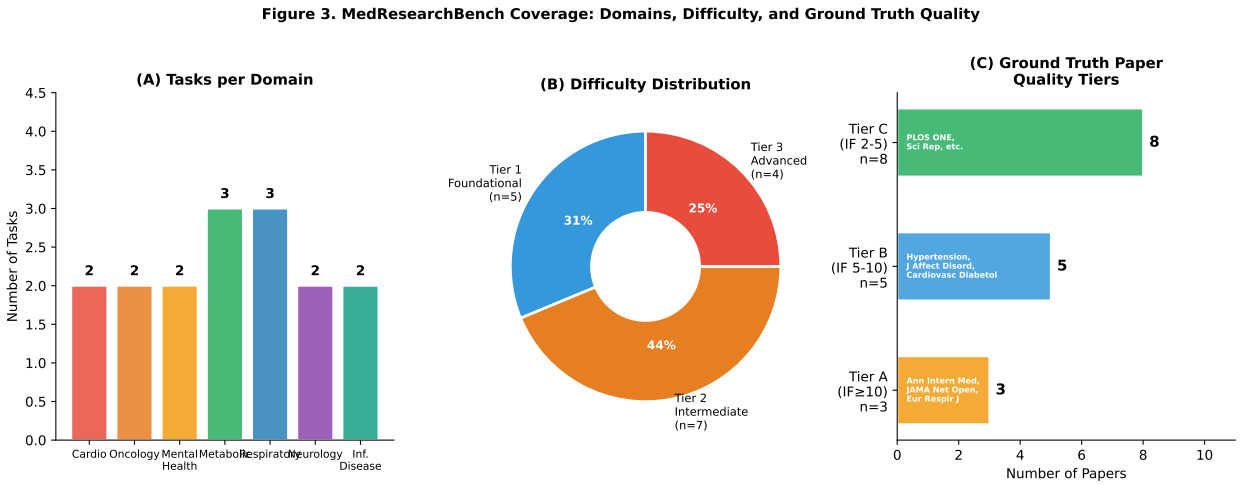

Figure 2: MedResearchBench coverage overview. (A) Tasks per clinical domain. (B) Difficulty tier distribution: Tier 1 Foundational (31%), Tier 2 Intermediate (44%), Tier 3 Advanced (25%). (C) Ground truth paper quality tiers by Impact Factor: 50% Tier C (IF 2–5), 31% Tier B (IF 5–10), 19% Tier A (IF ≥10).

4.5 Dataset Utilization

Table 9: Dataset utilization across MedResearchBench tasks.

Dataset	# Tasks	Cycles/Years Used
NHANES (standard cycles)	11	1999–2000 through 2017–2018 (up to 9 cycles)
NHANES III (with NDI linkage)	2	1988–1994, mortality follow-up through 2019
SEER / SEER-Medicare	2	1991–2018
Multi-source	1	Respiratory_001 spans both NHANES eras

NHANES dominance (13/16 tasks) reflects both its public availability and its central role in U.S. population health research. The 2 SEER-based oncology tasks provide important methodological contrast (registry-based survival analysis vs. survey-weighted cross-sectional analysis).

4.6 Comparison with ResearchClawBench Task Characteristics

Table 10: Structural comparison between ResearchClawBench and MedResearchBench.

Feature	**ResearchClawBench**	**MedResearchBench**
Total tasks	41	16 (Phase 1)
Domains	10 (fundamental science)	8 (clinical medicine)
Data paradigm	Experimental/simulated, in-task	Population survey/registry, downloadable
Typical task output	Code + figures + report	Full manuscript (STROBE-compliant)
Ground truth	Published paper (PDF included)	Published paper (metadata + checklist)
Checklist items per task	Variable	5 (standardized)
Evaluation dimensions	General	6 medical-specific dimensions
Unique requirement	—	Codebook with survey design variables
Confounding control	Not evaluated	Dedicated dimension

The key structural difference is that MedResearchBench tasks require navigating real-world data complexity (missing data, complex sampling, confounding) rather than executing well-defined computational experiments.

4.7 Initial Baseline Evaluation

To validate benchmark executability and establish a quantitative baseline, we evaluated an agentic data2paper pipeline on 3 pilot tasks spanning all three difficulty tiers. The pipeline follows a complete data-to-manuscript workflow: NHANES data download, survey-weighted statistical analysis (R + `survey` package), figure generation, STROBE-format manuscript writing, and automated checklist scoring.

Table 11: Pilot baseline results (AI Research Army pipeline, n=3 tasks).

Task	**Domain**	**Tier**	**Score**	**Grade**
Cardio_000	Cardiovascular	1	72/100	B
Mental_000	Mental Health	2	69/100	C+
Metabolic_002	Metabolic	3	75/100	B
Mean			**72/100**	**B**

Key findings from the pilot evaluation:

1. **Survey-weighted methodology: 100% compliance.** All three tasks correctly implemented complex survey design (NHANES MEC weights, PSU, strata), satisfying the benchmark's core methodological requirement. The Metabolic_002 task additionally handled NHANES III fixed-width format data and NDI mortality linkage, achieving the highest methodology score (92/100).

Table 12: Dimension-level scores across pilot tasks.

Dimension	Weight	Cardio_000	Mental_000	Metabolic_002	Mean
statistical_methodology	0.25	60	78	**92**	76.7
results_accuracy	0.25	72	52	52	58.7
confounding_sensitivity	0.20	68	76	75	73.0
clinical_interpretation	0.15	**88**	86	76	83.3
reporting_compliance	0.15	80	62	**88**	76.7
Weighted Total	1.00	**72**	**69**	**75**	**72**

2. **Clinical interpretation is the strongest dimension** (mean 83.3/100). The pipeline consistently produced mechanistic explanations, acknowledged limitations, and drew clinically actionable conclusions.

3. **Results accuracy is the primary limitation** (mean 58.7/100). Effect sizes were systematically attenuated across all tasks relative to ground truth papers. Cardio_000 had an acceptable OR deviation; Mental_000 had a reference group misspecification (7h used, target 8h); Metabolic_002 had attenuated low-normal TSH HR (1.15 vs. target 1.39, −17%).

4. **Sample sizes consistently 7–26% smaller** than target studies, attributable to stricter covariate completeness requirements (complete-case analysis vs. indicator variable imputation).

5. **Score does not decrease with task complexity** (Tier 1: 72, Tier 2: 69, Tier 3: 75), suggesting the pipeline handles advanced methodologies (Cox PH, NHANES III data engineering) without additional penalty.

These pilot results demonstrate that the benchmark is executable end-to-end by a current agentic pipeline and produces discriminative scores across dimensions and tasks. Detailed evaluation reports are available in the repository (https://github.com/TerryFYL/MedResearchBench/tree/main/bench-runs).

5 Discussion

5.1 Why Medical Research Needs Its Own Benchmark

Medical clinical research is among the most consequential domains of scientific inquiry, directly influencing treatment guidelines, public health policy, and patient outcomes. Existing AI research benchmarks, while valuable, do not capture these stakes. A machine learning benchmark evaluates whether an AI can optimize an accuracy metric; MedResearchBench evaluates whether an AI can navigate the inferential, ethical, and clinical complexity of determining, for example, whether dietary sodium intake is causally related to hypertension (Cardio_000) or whether racial disparities in cancer staging persist despite expanded screening coverage (Onco_002).

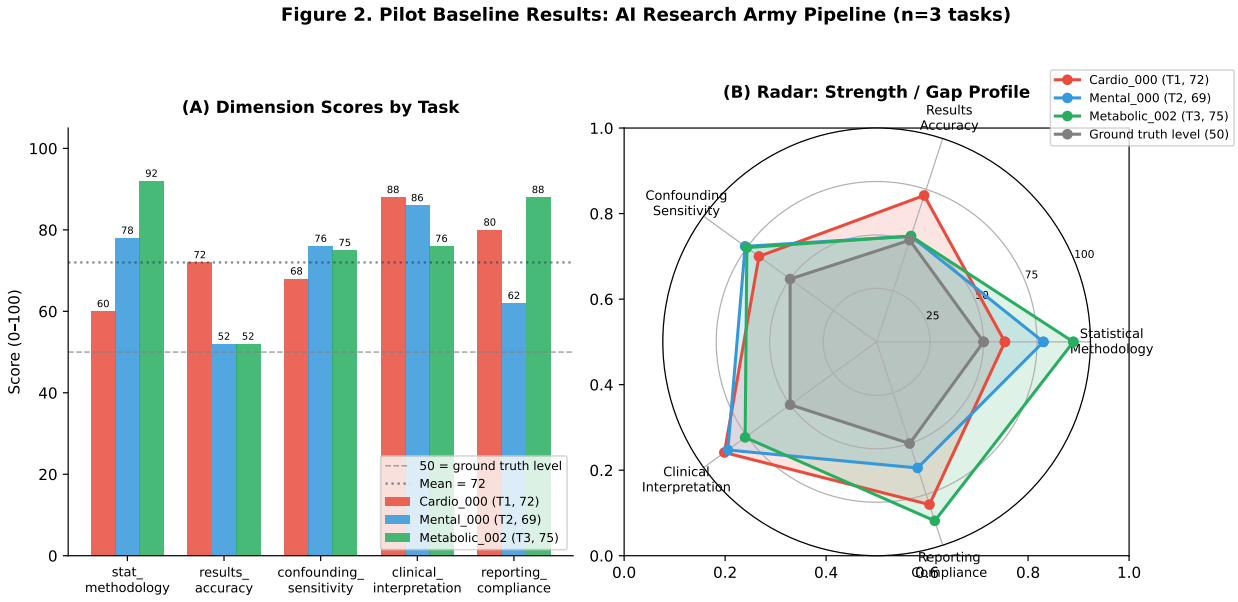

Figure 3: Pilot baseline results for the AI Research Army pipeline across 3 tasks. (A) Grouped bar chart of dimension-level scores; dashed line indicates ground truth level (50) and dotted line the mean score (72). (B) Radar chart showing the strength/gap profile: clinical interpretation is consistently the strongest dimension while results accuracy is the primary gap.

5.2 Comparison with ResearchClawBench

MedResearchBench is designed as a **complement** to ResearchClawBench, not a replacement. Together, the two benchmarks span the full landscape of AI research evaluation, from fundamental science to clinical medicine.

5.3 Responsible Use: Quality Gates Against Paper Mills

The NHANES paper mill problem (PLOS Biology, 2025) demonstrates that the combination of publicly available data and automated analytical pipelines can produce high-volume, low-quality publications that degrade the medical literature. MedResearchBench is designed to serve as a quality gate against this risk through four mechanisms: (1) the `confounding_sensitivity` dimension explicitly evaluates beyond naive single-variable associations; (2) the `clinical_interpretation` dimension requires clinically meaningful discussion; (3) the `reporting_compliance` dimension checks STROBE adherence; and (4) scoring against published papers of varying quality tiers enables identification of systems that produce paper-mill-quality output (scoring well below 50).

5.4 Limitations

1. **Phase 1 scope**: 16 tasks across 8 domains. Phase 2 will expand to 32 tasks with additional study designs.

2. **Dataset concentration**: 13/16 tasks use NHANES. Future phases will incorporate MIMIC, UK Biobank, and CMS public use files.

3. **Observational studies only**: Phase 1 excludes randomized controlled trials.

4. **LLM Judge limitations**: Automated evaluation may not capture all nuances of medical research quality. Expert physician review is planned for a validation subset.

5. **English-only**: All tasks and evaluation materials are in English.

6. **No wet-lab or clinical data collection**: Tasks evaluate the analytical pipeline only.

5.5 Future Directions

1. **Comparative system evaluation**: Having established a B-level (72/100) baseline with an agentic data2paper pipeline, future work will evaluate competing AI research systems (AI Scientist, data-to-paper, Agent Laboratory, single-LLM baselines) on the same tasks, populating the first public leaderboard for medical AI research quality.

2. **Phase 2 expansion**: 32 tasks covering all 8 domains with 4 tasks each, incorporating case-control studies, meta-analyses, and clinical trial analysis.

3. **Community contribution**: Open contribution framework enabling domain experts to add tasks following standardized templates.

4. **Clinical validation study**: Expert physician panel evaluation of a subset of AI-generated research outputs.

5. **Integration with regulatory frameworks**: Alignment with emerging guidelines for AI-generated research (SPIRIT-AI, CONSORT-AI extensions).

6 Conclusion

We present MedResearchBench, the first benchmark specifically designed for evaluating AI systems on medical clinical research tasks. The benchmark covers 8 clinical domains with 16 carefully curated tasks, introduces 6 medical-specific evaluation dimensions, and calibrates scoring against high-quality published papers using publicly available datasets. Our analysis of task characteristics reveals substantial methodological diversity (23 distinct statistical methods), a deliberate quality-tier distribution in ground truth papers (IF 2.3–51.0), and a three-tier difficulty stratification enabling progressive system evaluation. Initial baseline experiments on 3 pilot tasks—spanning all three difficulty tiers—demonstrate benchmark executability and yield a mean score of 72/100 (B-level) for an agentic data2paper pipeline. Survey-weighted methodology was correctly implemented across all tasks; the primary performance gap lies in results accuracy (mean 58.7/100), driven by covariate incompleteness and reference group misspecification. As AI systems increasingly automate scientific research, the medical domain—with its high stakes, complex methodology, and established quality standards—demands specialized evaluation. MedResearchBench provides that evaluation

platform and, through its anti-paper-mill design, serves as a quality gate for responsible AI-assisted medical research.

References

1. Lu C, Lu C, Lange RT, et al. The AI Scientist: Towards Fully Automated Open-Ended Scientific Discovery. *Nature*. 2024.

2. Ifargan S, Hafner M, Kern S, et al. Autonomous LLM-driven research from data to human-verifiable research papers. *NEJM AI*. 2024.

3. Schmidgall S, et al. Agent Laboratory: Using LLM Agents as Research Assistants. *EMNLP*. 2025.

4. Schmidgall S, et al. ResearchClawBench: A Benchmark for AI Research Agents. 2025.

5. Huang J, et al. AI-Researcher: Automating Scientific Research with Autonomous AI Agents. *NeurIPS Spotlight*. 2025.

6. Swanson K, et al. Virtual Lab: AI Agents Design New Nanobodies with Experimental Validation. *Nature*. 2024.

7. Beel J, et al. Robin: A Research AI Agent System. FutureHouse. 2025.

8. Byrne J, Labbé C, et al. The NHANES/CHARLS Paper Mill. *PLOS Biology*. 2025.

9. von Elm E, Altman DG, Egger M, et al. The Strengthening the Reporting of Observational Studies in Epidemiology (STROBE) Statement. *Ann Intern Med*. 2007;147(8):573–577.

10. Aggarwal R, Chiu N, Wadhera RK, et al. Racial/Ethnic Disparities in Hypertension. *Hypertension*. 2021;78(6):1719–1726.

11. Dong L, Xie Y, Zou X. Association between sleep duration and depression in US adults. *J Affect Disord*. 2022;296:120–126.

12. Yu B, Zhang X, Wang C, et al. Trends in depression among Adults in the United States, NHANES 2005-2016. *J Affect Disord*. 2020;263:609–620.

13. Pankratz VS, et al. Colorectal Cancer Survival Trends by Race and Stage. *Cancer Control*. 2022;29.

14. Liang PS, et al. Trends in Sociodemographic Disparities in CRC Staging and Survival. *Clin Transl Gastroenterol*. 2020;11(3):e00155.

15. Le MH, et al. Prevalence of HBV Vaccination Coverage, 1999 to 2016. *JAMA Network Open*. 2020;3(11):e2022388.

16. Rosenblum HG, et al. HPV Vaccine Impact Through 12 Years. *Ann Intern Med.* 2022;175(7):918–926.

17. Guess J, Beltran TH, Choi YS. Prediction of Metabolic Syndrome Using HOMA-IR. *Metab Syndr Relat Disord.* 2023;21(3):152–157.

18. Wang T, et al. NAFLD Prevalence and Advanced Fibrosis, NHANES 2011-2018. *Ann Hepatol.* 2024;29(1):101162.

19. Lu K, et al. Serum metals and cognitive performance in older adults. *Br J Nutr.* 2023;130(6):1040–1049.

20. Huang F, et al. Dietary antioxidant index and Parkinson's disease. *Front Aging Neurosci.* 2025;17:1522593.

21. Mendy A, et al. Blood VOCs and Airflow Obstruction. *Eur Respir J.* 2022;61(1):2200468.

22. Martinez CH, et al. Undiagnosed Obstructive Lung Disease in the US. *Ann Am Thorac Soc.* 2015;12(12):1788–1795.

23. Zhu P, Lao G, et al. TSH and mortality in euthyroid diabetic adults. *Cardiovasc Diabetol.* 2024.

24. Yin H, Huang W, Yang B. METS-IR and obstructive sleep apnea. *Sci Rep.* 2025.

25. Luo X, et al. Dietary sodium and hypertension, NHANES 2017-2018. *PLOS ONE.* 2026.

A Complete Task List with Metadata

#	Task ID	Domain	Research Question	Tier	Dataset
1	Cardio_000	Cardiovascular	Dietary sodium and hypertension dose-response	1	NHANES 2017-2018
2	Cardio_003	Cardiovascular	Racial disparities in hypertension control	2	NHANES 2013-2018
3	Onco_000	Oncology	CRC survival trends by race/ethnicity	2	SEER 1992-2018
4	Onco_002	Oncology	SES disparities in CRC staging and survival	3	SEER-Medicare
5	Mental_000	Mental Health	Sleep duration and depression, U-shaped	2	NHANES 2009-2016
6	Mental_002	Mental Health	Depression prevalence trends 2005-2016	1	NHANES 2005-2016
7	Metabolic_000	Metabolic	HOMA-IR threshold for metabolic syndrome	1	NHANES 2011-2018
8	Metabolic_002	Metabolic	TSH and all-cause/CV mortality (nonlinear)	3	NHANES III + NDI
9	Metabolic_003	Metabolic	NAFLD prevalence and fibrosis risk factors	1	NHANES 2011-2018

#	Task ID	Domain	Research Question	Tier	Dataset
10	Neuro_002	Neurology	Heavy metal exposure and cognitive function	2	NHANES 2011-2014
11	Neuro_003	Neurology	Dietary antioxidants and PD mortality	3	NHANES + NDI
12	Infection_000	Infectious Disease	HBV seroprevalence and vaccine trends	2	NHANES 1999-2016
13	Infection_001	Infectious Disease	HPV vaccine effectiveness (12-year)	3	NHANES 2003-2018
14	Respiratory_000	Respiratory	Air pollution biomarkers and lung function	1	NHANES 2007-2012
15	Respiratory_001	Respiratory	Undiagnosed COPD prevalence and mortality	3	NHANES III + 2007-2012
16	Respiratory_003	Respiratory	OSA risk and metabolic comorbidities	2	NHANES 2005-2018

B LLM Judge Prompt Template

The following prompt template is provided to the LLM Judge (GPT-4o or Claude Opus) for each task evaluation.

```
You are an expert medical research reviewer evaluating an AI-
   generated
research report against a published ground truth paper.

## Task Context
Research question: {task_description}
Dataset: {dataset_source}
Study design: {study_design}

## Ground Truth Checklist
{checklist_items_json}

## AI-Generated Report
{ai_report_text}

## Scoring Instructions
For each checklist item, assign a score from 0 to 100:
- 0: Completely absent or fundamentally wrong
- 25: Attempted but with major errors or omissions
- 50: Matches the quality of the ground truth published paper
- 75: Exceeds the ground truth paper in rigor or completeness
- 100: Exceptional quality, substantially beyond the published
   paper

[Apply dimension-specific criteria for statistical_methodology,
results_accuracy, visualization_quality,
   clinical_interpretation,
confounding_sensitivity, reporting_compliance]
```

```
## Output Format
For each checklist item: Item ID, Score (0-100), Brief
   justification
Then: Weighted total score + Overall qualitative assessment
```

C Paper Selection Criteria

Inclusion Criteria: (1) Published in a peer-reviewed journal with IF $\geq$ 2.0; (2) Uses one of the designated publicly available datasets; (3) Addresses the specific research question; (4) Employs appropriate statistical methods; (5) Available via PubMed with valid PMID and DOI.

Exclusion Criteria (Anti-Paper-Mill Filtering): (1) Formulaic single-factor association studies without analytical innovation; (2) Papers from known paper mill publishers; (3) Studies with inadequate confounding control (<3 covariates adjusted); (4) Papers not following minimum STROBE reporting standards; (5) Retracted or corrected papers.